

- Per-agent credit for multi-agent flutter suppression is derived in closed form.
- Closed-form credit beats naive and value-decomposition baselines after correction.
- Exploration scale, not algorithm choice, dominated performance by a factor of two.
- Six-rung classical baseline ladder and UVLM cross-validation ground the comparison.
- Learned policies trail decentralised LQG on average, but hold a condition it cannot.